

"WHAT LIES BEHIND YOU AND WHAT LIES IN FRONT OF  
YOU, PALES IN COMPARISON TO WHAT LIES INSIDE OF YOU."  
— **RALPH WALDO EMERSON**





# BAYESIAN STATISTICS

CLASS 2

# WHAT DID WE LEARN FROM CLASS 1?

- Terminology: Prior, Sampling distribution, Posterior
- How to define problem (decide sampling distribution of data, define priors for parameters, use pymc to generate posterior distribution of parameters)
- How to use posterior to answer questions about the parameter
- How data (sample size) and prior contribute to the posterior
- Why prior is VERY important when sample size is small



# GOALS FOR TODAY

- MCMC - Markov Chain Monte Carlo
  - What it is
  - Has it converged
  - Options to help convergence
- Options in running MCMC to get posterior distribution
- Another in-class example





# SIMULATING A DISTRIBUTION

```
y=rnorm(2000)  
> hist(y)  
> plot(y,type='l')
```





# ANOTHER DISTRIBUTION

```
> y=rbinom(2000,100,0.3)  
> hist(y)  
> plot(y,type='l')
```



# SKEWED DISTRIBUTION

```
> y=rbeta(2000,2,50)  
> hist(y)  
> plot(y,type='l')
```




# MARKOV CHAIN MONTE CARLO







# MARKOV CHAIN MONTE CARLO




# MARKOV CHAIN MONTE CARLO



# MARKOV CHAIN MONTE CARLO




# MARKOV CHAIN MONTE CARLO

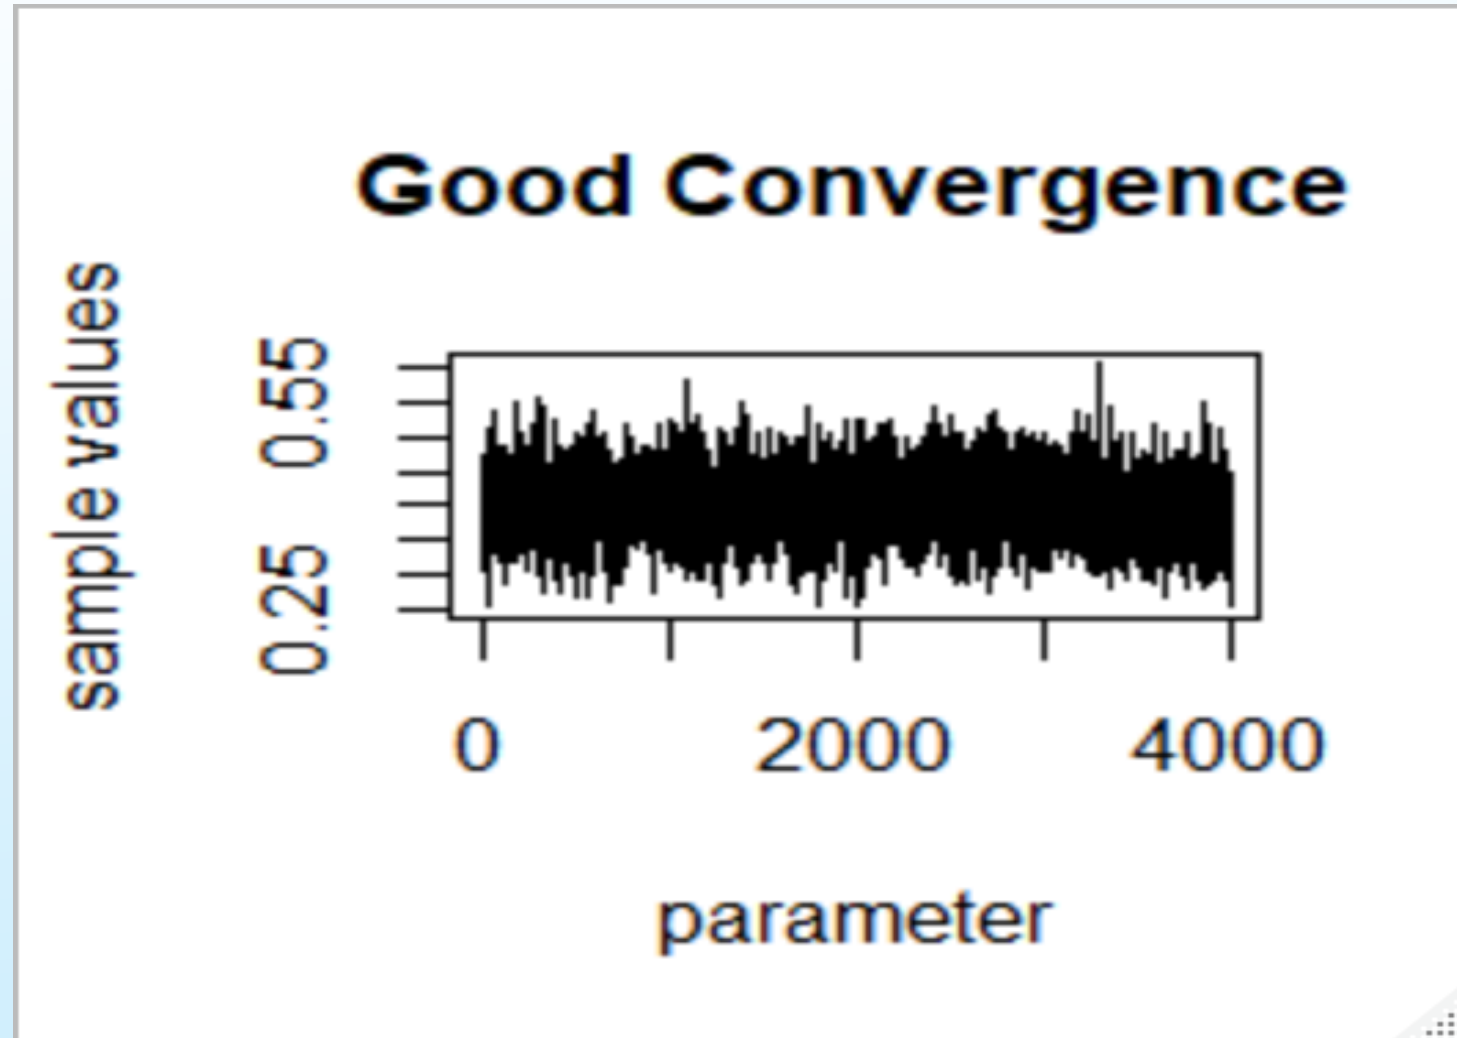


pymc uses the Hamiltonian Monte Carlo method for its Markov Chain and its adaptive variant called the no U-turn sampler (NUTS). For an interesting read, see <https://towardsdatascience.com/python-hamiltonian-monte-carlo-from-scratch-955dba96a42d/>  
Stan also uses NUTS.

# MCMC


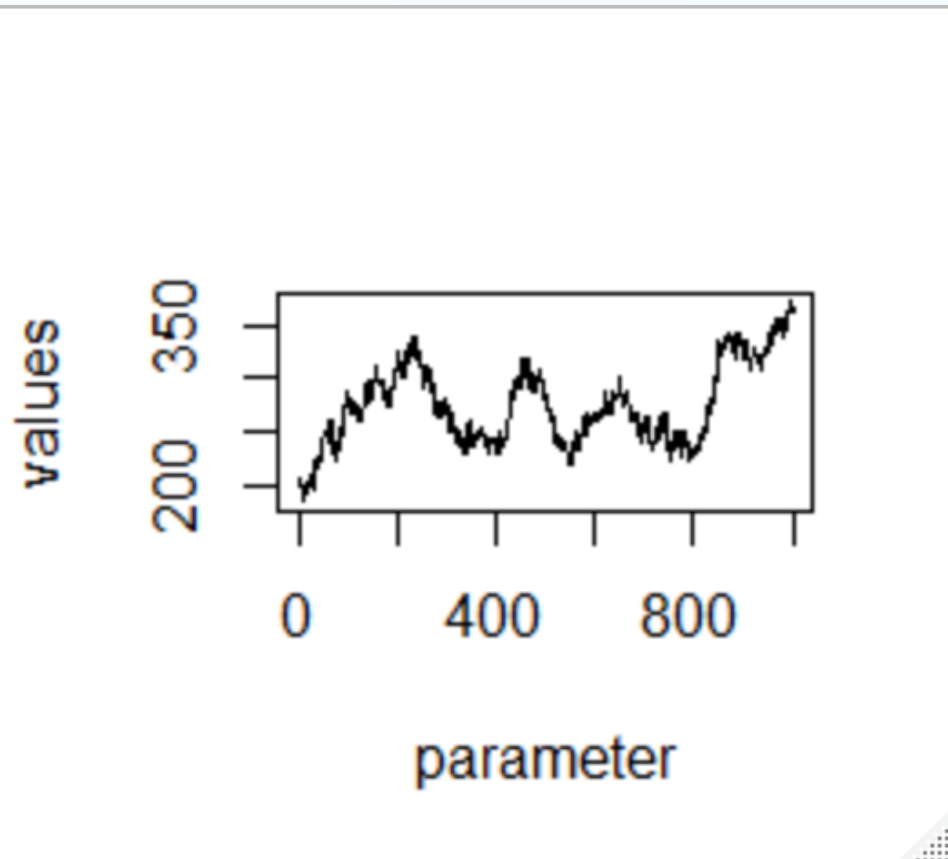


# CONVERGENCE





# NONCONVERGENCE



# FIXES

- Improper posterior or bad prior
  - Fix: New prior distribution
- Hasn't converged yet
  - Let the chain run longer
- Too much autocorrelation in chain
  - Thin the chain

# OPTIONS FOR PYMC CODE

```
n = 100
y = 40
# Model
with pm.Model () as binom_model :
    p = pm.Beta ("p", alpha=1, beta=1) ### Prior distribution
    y_obs = pm.Binomial("y_obs", n=n, p=p, observed=y) ### Sampling distribution
    trace = pm.sample (draws=3000, # samples after burn-in
tune=3000, # Warmup iterations
chains=4, # Number of Markov chains
thin=3, # Thinning factor (really doesn't work)
    random_seed=98763,
    return_inferencedata=True ) ### specify information
az.summary(trace) ### Summary
p1 = trace.posterior["p"].values.flatten()
thin_val=3
p_post = p1[:,::thin_val] ### would use this for inference
```

Creates four chains; each chain has 3000 values (it automatically throws away burn-in), thin does not work.... Need to thin after you pull off values (if autocorrelation is an issue)

p1 has 12000 observations (3000x4) and p\_post has 4000 observations (after thinning)

# TRACEPLOT OF FOUR CHAINS

## Code

```
import seaborn as sns

df_p = az.extract(trace, var_names=["p"]).to_dataframe()

sns.set(style="whitegrid")

plt.figure(figsize=(10, 6))

sns.lineplot(data=df_p, x="draw", y="p", hue="chain",
             alpha=0.8)

plt.xlabel("Draw")

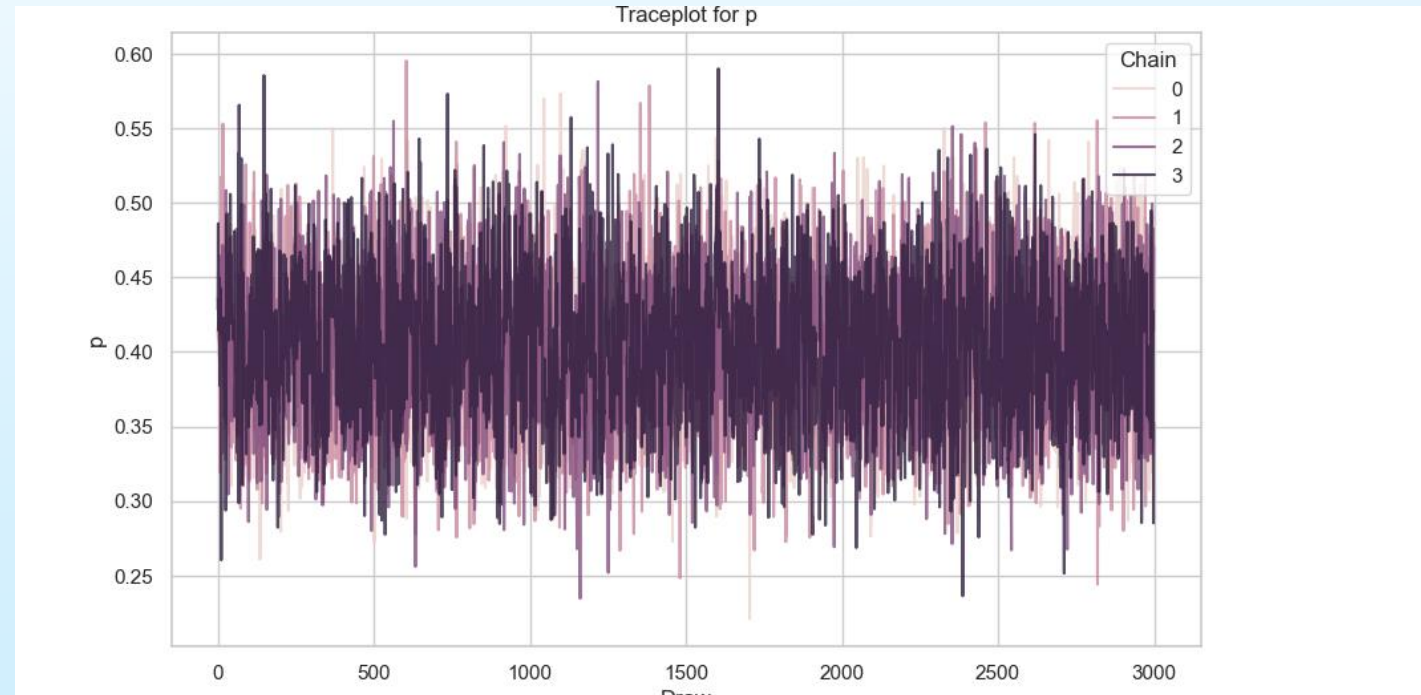
plt.ylabel("p")

plt.title("Traceplot for p")

plt.legend(title="Chain")

plt.show()
```

## Traceplot



# DIAGNOSTICS

```
az.summary(trace)
```

	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_bulk	ess_tail	r_hat
p	0.403	0.048	0.313	0.491	0.001	0.0	5323.0	8853.0	1.0

Want this to be close to 1 (this means convergence); if greater than 1.1, then there could potentially be a problem

# DIAGNOSTICS

```
az.summary(trace)
```

	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_bulk	ess_tail	r_hat
p	0.403	0.048	0.313	0.491	0.001	0.0	5323.0	8853.0	1.0

Effective sample size...bigger is better. If either of these are small, then there is a potential issue with independence of samples (ess\_bulk looks at the middle range of your posterior values while ess\_tail looks at the tails of the distribution). FIX if these are small: consider thinning.



# ANOTHER POTENTIAL WARNING

Divergences after tuning: if your chains are diverging after tuning, you can try to increase `target_accept`, increase burn-in period or try changing prior distributions.

VALUE AT RISK (VAR)

# ESTIMATE VALUE AT RISK

- Recall from Simulation and Risk calculating Value at Risk (VaR)
- Say we are interested in estimating  $\text{VaR}_{0.01}$  for Apple stock (AAPL) in Rate of Change (ROC) for one day
- If we assume Rate of Change ( $R_t$ ) follows a Normal distribution with mean  $\mu$  and standard deviation  $\sigma$
- We then need to assign a distribution to  $\mu$  and  $\sigma^2$ 
  - Assume  $\mu$  is distributed as  $\text{Normal}(0,100)$
  - Assume  $\sigma^2$  is distributed as  $\text{Inv-Gamma}(0.001,0.001)$
- Once we get posterior for  $\mu$  and  $\sigma^2$ , we can use this to get the 1<sup>st</sup> quantile

```
import pandas as pd
import yfinance as yf
import math

# Define the ticker symbol
ticker = 'AAPL'



# Download stock data
data = yf.download(ticker, start='2024-01-01')

# Calculate returns
data['aapl_r'] = data['Close'].pct_change()

# Select last 200 rows of Close and aapl_r columns
stocks = data[['Close', 'aapl_r']].tail(200)
```

```
import scipy.stats as stats
# Define the data
n = len(stocks)
y = stocks['aapl_r'].values
```

```
# Define the PyMC model
with pm.Model() as model:
    # Priors
    mu = pm.Normal('mu', mu=0, sigma=100)
    sigma2 = pm.InverseGamma('sigma2', alpha=0.001, beta=0.001)
    # Likelihood
    y_lik = pm.Normal('y_lik', mu=mu, sigma=pm.math.sqrt(sigma2),
                      observed=y)
    # Sample posterior
    trace = pm.sample(1000, tune=500)
```



Looks like we have convergence

```
az.summary(trace)
```

```
]:
```

	mean	sd	hdi_3%	hdi_97%	mcse_mean	mcse_sd	ess_bulk	ess_tail	r_hat
<b>mu</b>	0.001	0.001	-0.001	0.003	0.0	0.0	4486.0	2875.0	1.0
<b>sigma2</b>	0.000	0.000	0.000	0.000	0.0	0.0	4080.0	3166.0	1.0




# USING POSTERIOR INFORMATION:

Posterior samples

[31]:

	<b>mu</b>	<b>sigma2</b>
<b>0</b>	-0.001192	0.000179
<b>1</b>	-0.000943	0.000188
<b>2</b>	-0.001685	0.000180
<b>3</b>	-0.001862	0.000179
<b>4</b>	-0.002084	0.000178



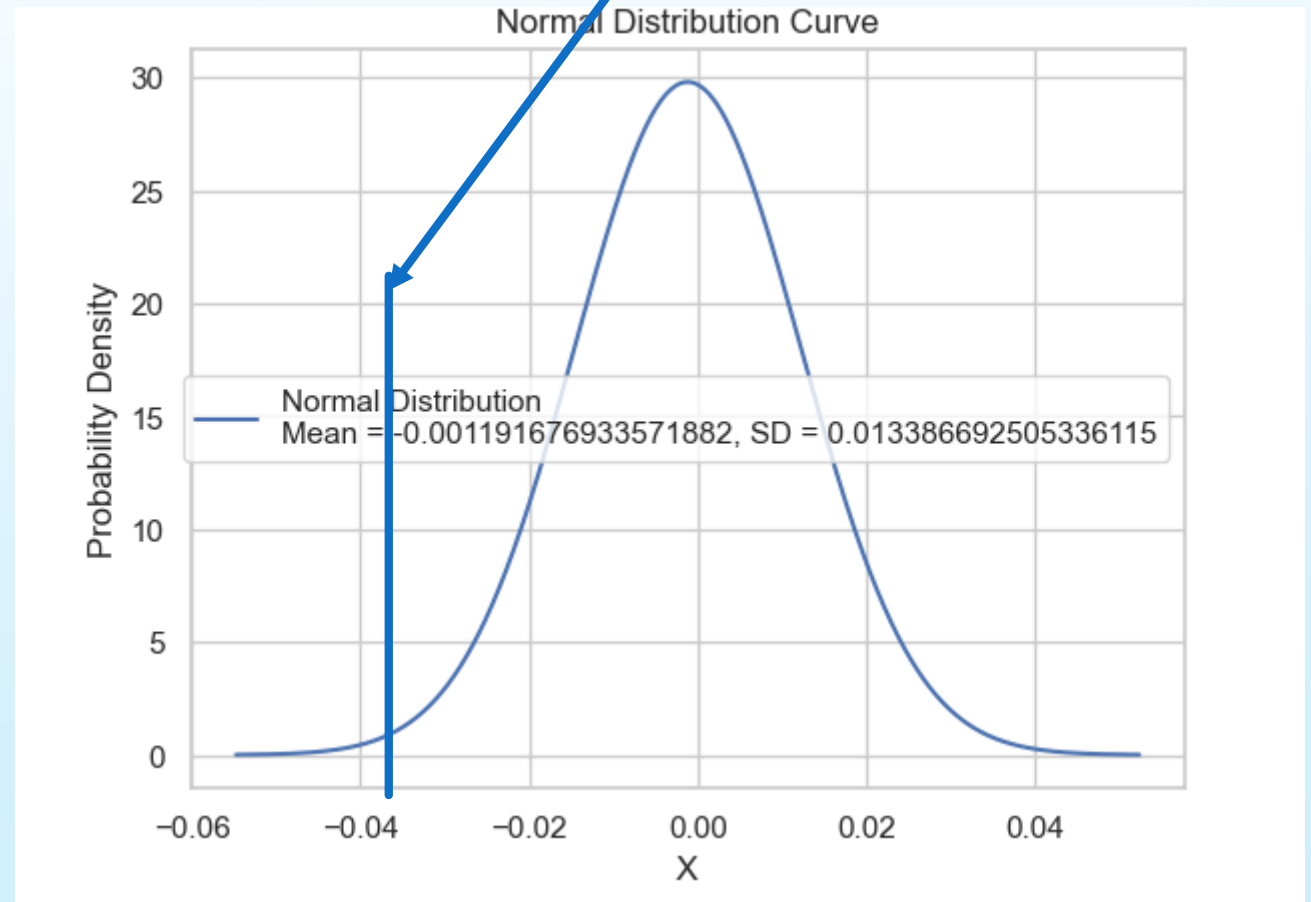
First quantile (VaR0.01)

# USING POSTERIOR INFORMATION:

Posterior samples

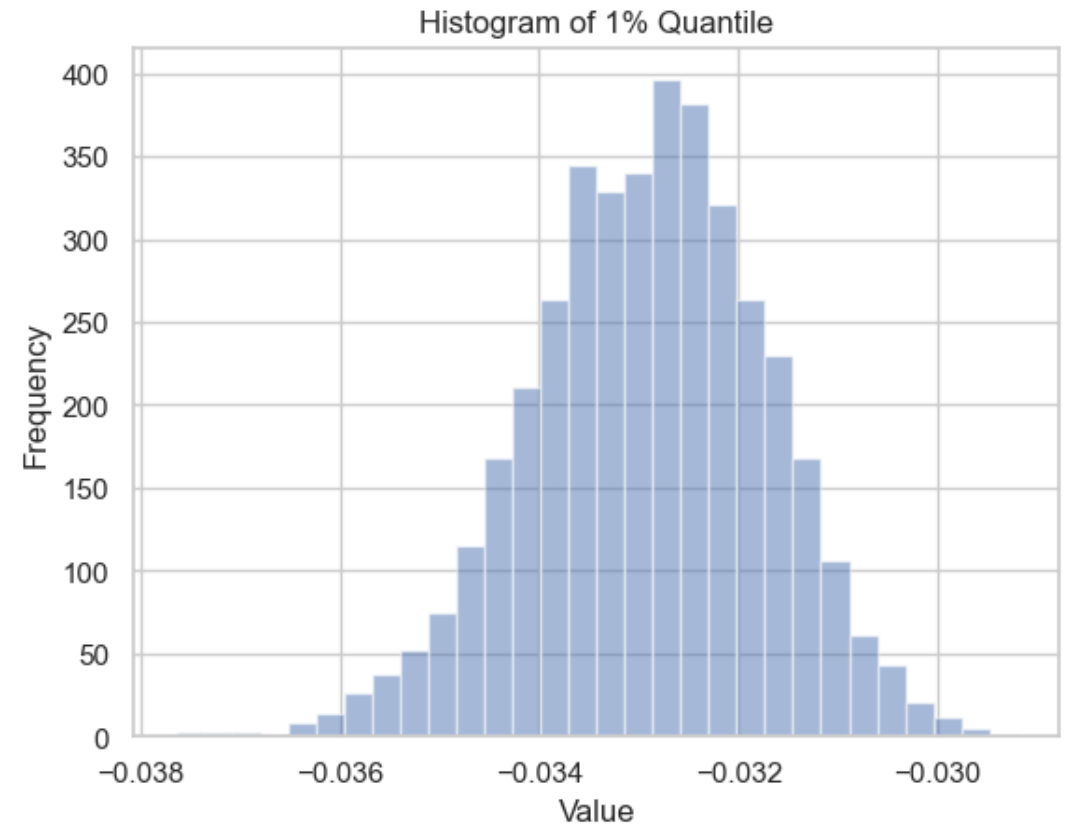
[31]:

	<b>mu</b>	<b>sigma2</b>
<b>0</b>	-0.001192	0.000179
<b>1</b>	-0.000943	0.000188
<b>2</b>	-0.001685	0.000180
<b>3</b>	-0.001862	0.000179
<b>4</b>	-0.002084	0.000178



## Posterior distribution of VaR0.01

```
posterior_mu = trace.posterior["mu"].values.flatten()
posterior_sigma2= trace.posterior["sigma2"].values.flatten()
# Calculate 1% quantile
one_per = stats.norm.ppf(0.01, loc=posterior_mu,
scale=np.sqrt(posterior_sigma2))
# Plot histogram of 1% quantile
plt.hist(one_per, bins=30, density=False, alpha=0.5)
plt.xlabel('Value')
plt.ylabel('Frequency')
plt.title('Histogram of 1% Quantile')
plt.grid(True)
plt.show()
```



**USE THIS INFO FOR IN-CLASS ASSIGNMENT**